

## 6. Object Detection using Transfer Learning of CNN Architectures

**Aim:** To implement object detection using transfer learning by leveraging a pre-trained Convolutional Neural Network (CNN) architecture.

**Theory:** **Object detection** is a computer vision task that involves identifying and localizing objects within an image by drawing bounding boxes around them and assigning a class label. Training an object detection model from scratch is computationally expensive and requires massive datasets. **Transfer learning** offers a highly effective solution.

The process leverages a CNN pre-trained on a large-scale image classification dataset (like ImageNet). This pre-trained model has already learned to extract rich hierarchical features from images.

**Steps:**

- **a. Load in a Pre-trained CNN Model:** We start by loading a popular CNN architecture (e.g., VGG16, ResNet50, MobileNetV2) that has been pre-trained on ImageNet. Typically, we load the **convolutional base** of the model, excluding its final classification layers, as these are task-specific.
- **b. Freeze Parameters (Weights) in the Model's Lower Convolutional Layers:** To retain the general feature extraction capabilities learned from ImageNet, we freeze the weights of the pre-trained convolutional base. This prevents them from being updated during the initial stages of training on our specific dataset, avoiding catastrophic forgetting of learned features.
- **c. Add a Custom Classifier with Trainable Parameters:** On top of the frozen convolutional base, we add new layers designed for object detection. This often includes:
  - A few Dense layers for high-level feature processing.
  - Specific layers tailored for object detection, such as region proposal networks or bounding box regression and classification heads (depending on the object detection framework used, e.g., Faster R-CNN, YOLO, SSD). For simplicity here, we'll illustrate adding a classifier for bounding box prediction and class probabilities.
  - **Dropout** layers can be added for regularization.
- **d. Train Classifier Layers on Training Data:** The model is compiled and trained. During this phase, only the weights of the newly added custom layers are updated. The pre-trained base acts as a fixed feature extractor. We use a suitable optimizer (e.g., Adam) and loss function. For object detection, the loss is typically a combination of classification loss (e.g., categorical cross-entropy) and regression loss (e.g., smooth L1 loss for bounding box coordinates).

- **e. Fine-tune Hyperparameters and Unfreeze More Layers (Optional but Recommended):** After the custom layers have been trained, we can optionally **fine-tune** the model. This involves:
  - **Unfreezing** some of the top layers of the pre-trained convolutional base.
  - Recompiling the model with a very **low learning rate**.
  - Training the model again. This allows the pre-trained features to adapt slightly to the specifics of the new task, potentially improving performance further. Hyperparameters like learning rate, batch size, and optimizer choice are critical here.

**Code Example (Conceptual - Fine-tuning for classification as a proxy for detection steps):**

*Note: Implementing a full object detection model (like Faster R-CNN or YOLO) is complex and beyond a simple manual example. This code demonstrates the core transfer learning concept by fine-tuning a pre-trained model for classification on a dataset like CIFAR-10, which is a stepping stone. For actual object detection, you'd integrate specific detection heads.*

Python

```
import tensorflow as tf

from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Flatten, Dropout
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers import Adam
import matplotlib.pyplot as plt

# --- a. Load in a pre-trained CNN model trained on a large dataset ---

# Load VGG16 pre-trained on ImageNet, without the top classification layers
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))
# input_shape is set to 32x32 for CIFAR-10. VGG16 was originally trained on 224x224,
# but Keras allows adapting it. However, for best results, consider resizing CIFAR-10 to
# 224x224 or using models designed for smaller inputs.
```

```
# --- b. Freeze parameters (weights) in the model's lower convolutional layers ---

# Freeze all layers in the base model initially
for layer in base_model.layers:
    layer.trainable = False

# --- c. Add a custom classifier with several layers of trainable parameters ---

x = base_model.output
x = Flatten()(x) # Flatten the output of the convolutional base
x = Dense(512, activation='relu')(x) # Add a dense layer
x = Dropout(0.5)(x) # Add dropout for regularization
# Add the final output layer for classification (CIFAR-10 has 10 classes)
num_classes = 10
predictions = Dense(num_classes, activation='softmax')(x) # Softmax for multi-class
classification

# Create the new model
model = Model(inputs=base_model.input, outputs=predictions)

model.summary()

# --- d. Train classifier layers on training data available for the task ---

# Load and preprocess CIFAR-10 data (similar to previous example)
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
y_train = to_categorical(y_train, num_classes)
y_test = to_categorical(y_test, num_classes)
```

```

# Compile the model

# Use Adam optimizer with a moderate learning rate for the initial training
model.compile(optimizer=Adam(learning_rate=0.001),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train the model (only the new layers are trained)
print("\nTraining the custom classifier layers...")
batch_size = 64
epochs = 5 # Reduced epochs for demonstration
history = model.fit(x_train, y_train,
                     batch_size=batch_size,
                     epochs=epochs,
                     validation_data=(x_test, y_test))

# --- e. Fine-tune hyperparameters and unfreeze more layers as needed ---
print("\nStarting fine-tuning...")

# Unfreeze some top layers of the base model

# For VGG16, let's unfreeze the last convolutional block (e.g., conv5)
# This requires knowing the layer names/indices. VGG16 has blocks 'block5_conv1',
# 'block5_conv2', 'block5_conv3'.
# Let's unfreeze from 'block5_conv1' onwards.

for layer in base_model.layers[20:]: # Layers from index 20 onwards in VGG16 are part of
# 'block5'
    layer.trainable = True

# Re-compile the model with a very low learning rate for fine-tuning
# A low learning rate prevents large updates that could destroy learned features

```

```
model.compile(optimizer=Adam(learning_rate=0.00001), # Very low learning rate
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary() # Check which layers are now trainable

# Train again for fine-tuning
print("Fine-tuning the model with unfreezed layers...")
fine_tune_epochs = 5 # More epochs for fine-tuning might be needed
total_epochs = epochs + fine_tune_epochs

history_fine = model.fit(x_train, y_train,
                         batch_size=batch_size,
                         epochs=fine_tune_epochs, # Train for additional epochs
                         initial_epoch=history.epoch[-1] + 1, # Continue from where we left off
                         validation_data=(x_test, y_test))

# Combine histories for plotting
history.history['accuracy'].extend(history_fine.history['accuracy'])
history.history['val_accuracy'].extend(history_fine.history['val_accuracy'])
history.history['loss'].extend(history_fine.history['loss'])
history.history['val_loss'].extend(history_fine.history['val_loss'])

# Plot training history after fine-tuning
plt.figure(figsize=(14, 5))

# Plotting Accuracy
plt.subplot(1, 2, 1)
plt.plot(range(1, total_epochs + 1), history.history['accuracy'], label='Training Accuracy')
```

```

plt.plot(range(1, total_epochs + 1), history.history['val_accuracy'], label='Validation Accuracy')

plt.title('Training and Validation Accuracy (After Fine-tuning)')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

# Plotting Loss
plt.subplot(1, 2, 2)
plt.plot(range(1, total_epochs + 1), history.history['loss'], label='Training Loss')
plt.plot(range(1, total_epochs + 1), history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss (After Fine-tuning)')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.show()

# Evaluate the final model
print("\nEvaluating the final fine-tuned model...")
loss, accuracy = model.evaluate(x_test, y_test, verbose=0)
print(f'Final Test Loss: {loss:.4f}')
print(f'Final Test Accuracy: {accuracy:.4f}')

Conclusion: Transfer learning significantly accelerates the development of effective models, especially for tasks like object detection where data is scarce or training is resource-intensive. By leveraging pre-trained features, we can build powerful models with less data and computational effort. Fine-tuning further refines the model's performance by adapting pre-

```

trained features to the specific nuances of the target task. This approach is fundamental to many state-of-the-art computer vision applications.